

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load all sheet names
excel_path = "Customer_Churn_Data_Large.xlsx"
xls = pd.ExcelFile(excel_path)
sheet_names = xls.sheet_names # List of all sheet names
print("Available sheets:", sheet_names)

# Load each sheet into a dictionary of DataFrames
dataframes = {sheet: xls.parse(sheet) for sheet in sheet_names}

# Access example:
df_main = dataframes['Customer_Demographics'] # Replace with actual
sheet name
df_main.head()

```

Available sheets: ['Customer_Demographics', 'Transaction_History', 'Customer_Service', 'Online_Activity', 'Churn_Status']

	CustomerID	Age	Gender	MaritalStatus	IncomeLevel
0	1	62	M	Single	Low
1	2	65	M	Married	Low
2	3	18	M	Single	Low
3	4	21	M	Widowed	Low
4	5	21	M	Divorced	Medium

```

def show_all_df_info(dataframes):
    for name, df in dataframes.items():
        print(f"\n=== {name} ===")
        print(df.info())
show_all_df_info(dataframes)

```

```

=== Customer_Demographics ===
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   CustomerID      1000 non-null   int64
1   Age             1000 non-null   int64
2   Gender          1000 non-null   object
3   MaritalStatus   1000 non-null   object
4   IncomeLevel     1000 non-null   object
dtypes: int64(2), object(3)
memory usage: 39.2+ KB
None

```

```

=== Transaction_History ===
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5054 entries, 0 to 5053
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            5054 non-null   int64
1   TransactionID          5054 non-null   int64
2   TransactionDate        5054 non-null   datetime64[ns]
3   AmountSpent            5054 non-null   float64
4   ProductCategory       5054 non-null   object
dtypes: datetime64[ns](1), float64(1), int64(2), object(1)
memory usage: 197.6+ KB
None

```

```

=== Customer_Service ===
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1002 entries, 0 to 1001
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            1002 non-null   int64
1   InteractionID          1002 non-null   int64
2   InteractionDate        1002 non-null   datetime64[ns]
3   InteractionType        1002 non-null   object
4   ResolutionStatus      1002 non-null   object
dtypes: datetime64[ns](1), int64(2), object(2)
memory usage: 39.3+ KB
None

```

```

=== Online_Activity ===
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            1000 non-null   int64
1   LastLoginDate         1000 non-null   datetime64[ns]
2   LoginFrequency        1000 non-null   int64
3   ServiceUsage          1000 non-null   object
dtypes: datetime64[ns](1), int64(2), object(1)
memory usage: 31.4+ KB
None

```

```

=== Churn_Status ===
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  -

```

```
0    CustomerID    1000 non-null    int64
1    ChurnStatus    1000 non-null    int64
dtypes: int64(2)
memory usage: 15.8 KB
None
```

□ Join Strategy Explanation

To build a comprehensive dataset for churn prediction, we need to consolidate customer information from all available sources. Each of the provided sheets contributes a different dimension of behavioral, demographic, or outcome data. Here's the logic behind our join strategy:

1. Customer_Demographics

This sheet provides static customer attributes such as age, gender, marital status, and income level. It serves as our **primary base table**, as each entry represents a unique customer (CustomerID).

2. Transaction_History

Each customer can have **multiple transactions**, making this a one-to-many relationship. Therefore, we **aggregate** this table by CustomerID to extract useful features such as:

- Total amount spent
- Average transaction value
- Number of transactions

3. Customer_Service

Similar to transaction history, each customer can have **multiple interactions** with customer service. Thus, we also **aggregate** this table, capturing features such as:

- Number of interactions
- Most common resolution type

4. Online_Activity

This data contains one row per customer, showing their **most recent login, login frequency**, and **digital product usage**. It can be **joined directly** on CustomerID without aggregation.

5. Churn_Status

This is our **target variable**, indicating whether the customer has churned. Since it has a one-to-one mapping with CustomerID, it can also be **joined directly**.

By merging these tables using CustomerID as the key, we construct a unified dataset containing demographic, transactional, service, digital behavior, and churn outcome data—ideal for performing EDA and training a predictive model.

```
# 1. Aggregating Transaction_History
agg_trans =
dataframes['Transaction_History'].groupby('CustomerID').agg({
    'AmountSpent': ['sum', 'mean', 'count'],
}).reset_index()
agg_trans.columns = ['CustomerID', 'TotalSpent', 'AvgSpent',
'NumTransactions']

# 2. Aggregating Customer_Service
agg_service =
dataframes['Customer_Service'].groupby('CustomerID').agg({
    'InteractionID': 'count',
    'ResolutionStatus': lambda x: x.value_counts().idxmax() # most
common resolution
}).reset_index()
agg_service.columns = ['CustomerID', 'NumInteractions',
'MostCommonResolution']

# 3. Merging all
df_merged = dataframes['Customer_Demographics'] \
    .merge(agg_trans, on='CustomerID', how='left') \
    .merge(agg_service, on='CustomerID', how='left') \
    .merge(dataframes['Online_Activity'], on='CustomerID', how='left')
\
    .merge(dataframes['Churn_Status'], on='CustomerID', how='left')

df_merged.head()
```

	CustomerID	Age	Gender	MaritalStatus	IncomeLevel	TotalSpent
0	1	62	M	Single	Low	416.50
1	2	65	M	Married	Low	1547.42
2	3	18	M	Single	Low	1702.98
3	4	21	M	Widowed	Low	917.29
4	5	21	M	Divorced	Medium	2001.49

	NumTransactions	NumInteractions	MostCommonResolution	LastLoginDate
0	1	1.0	Resolved	2023-10-21
1	7	1.0	Resolved	2023-12-05

2	6	1.0	Resolved	2023-11-15
3	5	2.0	Resolved	2023-08-25
4	8	NaN	NaN	2023-10-27

	LoginFrequency	ServiceUsage	ChurnStatus
0	34	Mobile App	0
1	5	Website	1
2	3	Website	0
3	2	Website	0
4	41	Website	0

```
df_merged.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null	Count	Dtype
---	-----	-----	-----	-----
0	CustomerID	1000	non-null	int64
1	Age	1000	non-null	int64
2	Gender	1000	non-null	object
3	MaritalStatus	1000	non-null	object
4	IncomeLevel	1000	non-null	object
5	TotalSpent	1000	non-null	float64
6	AvgSpent	1000	non-null	float64
7	NumTransactions	1000	non-null	int64
8	NumInteractions	668	non-null	float64
9	MostCommonResolution	668	non-null	object
10	LastLoginDate	1000	non-null	datetime64[ns]
11	LoginFrequency	1000	non-null	int64
12	ServiceUsage	1000	non-null	object
13	ChurnStatus	1000	non-null	int64

```
dtypes: datetime64[ns](1), float64(3), int64(5), object(5)
```

```
memory usage: 109.5+ KB
```

Correlation Analysis

```
# Copy and convert categorical columns explicitly
```

```
df_corr = df_merged.copy()
```

```
# Map gender if it's still object dtype
```

```
if df_corr['Gender'].dtype == 'object':
```

```
    df_corr['Gender'] = df_corr['Gender'].map({'M': 0, 'F': 1})
```


```
# Also convert other categorical columns if you'd like to include them
```

```

df_corr['ServiceUsage'] =
df_corr['ServiceUsage'].astype('category').cat.codes
df_corr['MaritalStatus'] =
df_corr['MaritalStatus'].astype('category').cat.codes
df_corr['IncomeLevel'] =
df_corr['IncomeLevel'].astype('category').cat.codes

# Then plot again
plt.figure(figsize=(10, 8))
corr_matrix = df_corr.corr(numeric_only=True)
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix of Numerical Features')
plt.show()

```



```

# Copy and encode if needed
df_focus = df_merged.copy()
df_focus['Gender'] = df_focus['Gender'].map({'M': 0, 'F': 1})
df_focus['ServiceUsage'] =
df_focus['ServiceUsage'].astype('category').cat.codes
df_focus['MaritalStatus'] =
df_focus['MaritalStatus'].astype('category').cat.codes
df_focus['IncomeLevel'] =
df_focus['IncomeLevel'].astype('category').cat.codes

# Compute correlations
correlations = df_focus.corr(numeric_only=True)
['ChurnStatus'].drop('ChurnStatus')

# Plot
plt.figure(figsize=(8, 6))
sns.barplot(x=correlations.values, y=correlations.index,
palette='coolwarm')
plt.title('Correlation of Features with ChurnStatus')
plt.xlabel('Correlation Coefficient')
plt.ylabel('Feature')
plt.grid(True, axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

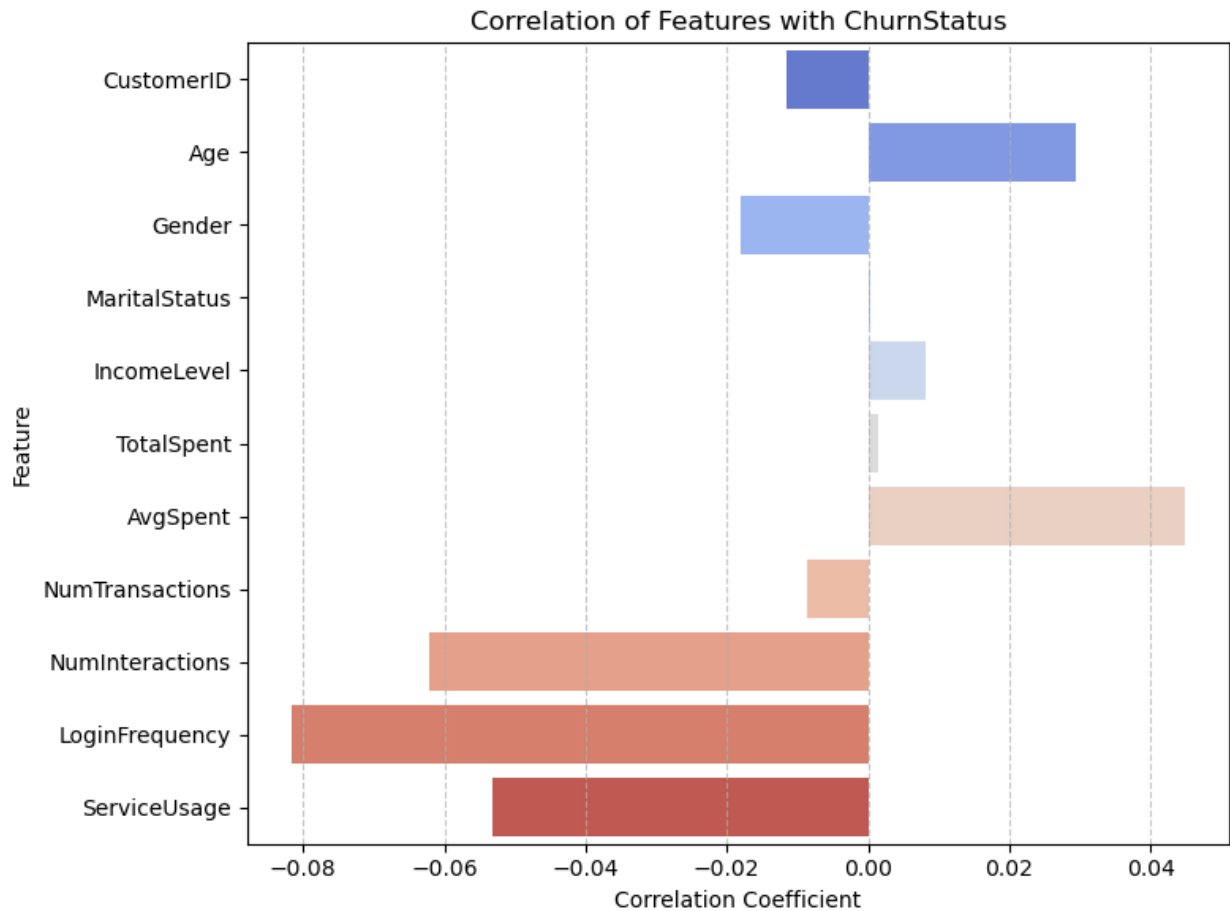
# Correlation with ChurnStatus (numerical output)
correlations = df_focus.corr(numeric_only=True)
['ChurnStatus'].drop('ChurnStatus')
correlations_sorted = correlations.sort_values(ascending=False)
print(correlations_sorted)

C:\Users\Kasandika Andariefli\AppData\Local\Temp\
ipykernel_4548\1141302335.py:13: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `y` variable to `hue` and set
`legend=False` for the same effect.

sns.barplot(x=correlations.values, y=correlations.index,
palette='coolwarm')

```



```
AvgSpent      0.044811
Age           0.029407
IncomeLevel   0.008134
TotalSpent    0.001324
MaritalStatus 0.000270
NumTransactions -0.008598
CustomerID    -0.011528
Gender        -0.018132
ServiceUsage  -0.053155
NumInteractions -0.062044
LoginFrequency -0.081615
Name: ChurnStatus, dtype: float64
```

□ Feature Correlation with ChurnStatus (Updated)

This visualization and summary highlight how each feature correlates with the target variable `ChurnStatus`, where a value of **1 represents a churned customer**. Understanding this directionality is key:

- **Positive correlation** → associated with **higher likelihood of churn**
- **Negative correlation** → associated with **lower likelihood of churn**

□ Interpretation:

- **LoginFrequency (-0.0816)** shows the strongest inverse correlation with churn, suggesting customers who log in frequently are more likely to **stay**.
- **NumInteractions (-0.0620)** and **ServiceUsage (-0.0532)** follow a similar trend—highlighting that **engagement and support usage reduce churn risk**.
- **Gender (-0.0181)**, now correctly encoded, suggests that **males (0)** may churn slightly more than females (1), though the effect is weak.
- **AvgSpent (+0.0448)** and **Age (+0.0294)** have small positive correlations, hinting that higher spending and older age might slightly increase churn—but this could indicate dissatisfaction or changing needs.
- **TotalSpent**, **IncomeLevel**, and **MaritalStatus** show near-zero correlation and likely require cross-feature interaction analysis for deeper insight.

□ Takeaway:

The top features negatively correlated with churn relate to **digital behavior and service usage**. These should be prioritized in model design and retention strategies. No single feature is dominant—supporting the use of a multivariate predictive model to capture subtle patterns.

Anomaly Detection

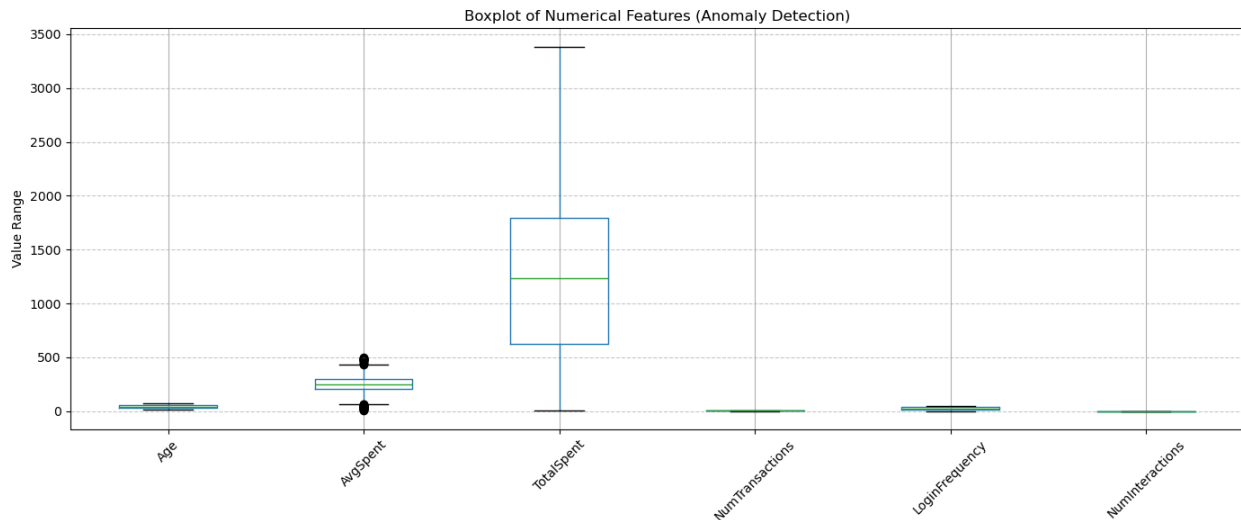
```
from scipy.stats import zscore

# Select numeric columns (excluding target)
numeric_cols = ['Age', 'AvgSpent', 'TotalSpent', 'NumTransactions',
                'LoginFrequency', 'NumInteractions']

plt.figure(figsize=(14, 6))
df_merged[numeric_cols].boxplot()
plt.title('Boxplot of Numerical Features (Anomaly Detection)')
plt.ylabel('Value Range')
plt.xticks(rotation=45)
plt.grid(True, axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Compute z-scores
z_scores = df_merged[numeric_cols].apply(zscore)

# Count outliers (Z-score > 3 or < -3)
outlier_counts = (z_scores.abs() > 3).sum()
print("Outliers (|z| > 3) per feature:")
print(outlier_counts)
```



```
Outliers ( $|z| > 3$ ) per feature:  
Age          0  
AvgSpent     6  
TotalSpent   0  
NumTransactions 0  
LoginFrequency 0  
NumInteractions 0  
dtype: int64
```

⚠ Anomaly Detection (Z-Score & Boxplot Analysis)

To detect unusual values that may indicate **data quality issues** or **rare behavioral patterns**, we applied **Z-score based anomaly detection** and visualized the results using boxplots.

📋 Z-Score Summary:

Only **AvgSpent** exhibits significant outliers — 6 customers had unusually high or low average spending patterns.

📊 Boxplot Insights:


- **AvgSpent** shows clear outliers and a fairly wide spread, suggesting variability in spending behavior. This may capture high-value or low-engagement customer extremes.
- **TotalSpent**, despite its large range, shows no statistical outliers by Z-score — possibly due to its spread being normalized across the population.
- Other features like **Age**, **LoginFrequency**, and **NumTransactions** are tightly distributed with no significant anomalies.

🔍 Interpretation:

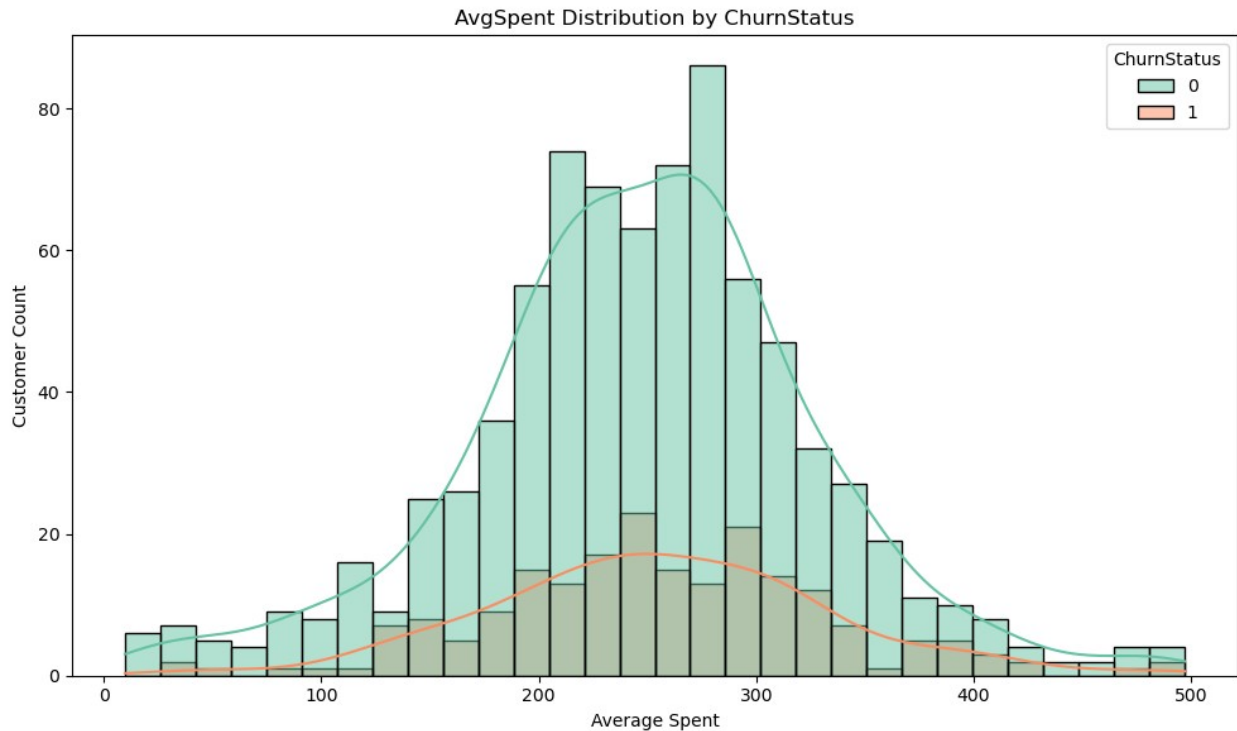
- The outliers in **AvgSpent** should be reviewed carefully:

- Are they data entry errors?
 - Or do they represent key behavioral patterns (e.g., high-value churn risks)?
- Since no other features show high-anomaly frequency, we **do not recommend removing these outliers at this stage**, but instead flagging them for potential stratified analysis or robust modeling.

```
plt.figure(figsize=(10, 6))
sns.histplot(df_merged['AvgSpent'], bins=30, kde=True,
color='steelblue')
plt.title('Distribution of Average Spending per Customer')
plt.xlabel('Average Spent')
plt.ylabel('Number of Customers')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
plt.figure(figsize=(10, 6))
sns.histplot(data=df_merged, x='AvgSpent', hue='ChurnStatus', bins=30,
kde=True, palette='Set2')
plt.title('AvgSpent Distribution by ChurnStatus')
plt.xlabel('Average Spent')
plt.ylabel('Customer Count')
plt.tight_layout()
plt.show()
```



```
high_spenders = df_merged[df_merged['AvgSpent'] >
df_merged['AvgSpent'].quantile(0.95)]
print("High spenders who churned:")
print(high_spenders[high_spenders['ChurnStatus'] == 1])
```

High spenders who churned:

	CustomerID	Age	Gender	MaritalStatus	IncomeLevel	TotalSpent
AvgSpent \						
136	137	21	F	Married	High	431.47
431.470000						
143	144	45	F	Married	Low	419.85
419.850000						
236	237	61	M	Divorced	Low	398.99
398.990000						
260	261	27	M	Married	Low	489.07
489.070000						
281	282	42	M	Single	High	2016.68
403.336000						
487	488	28	F	Married	Medium	393.84
393.840000						
497	498	45	F	Single	Low	763.91
381.955000						
509	510	30	M	Single	Medium	392.10
392.100000						
564	565	69	M	Single	Low	2436.25
406.041667						
613	614	64	F	Married	High	1239.16

413.053333						
629	630	61	F	Married	Medium	942.44
471.220000						
826	827	42	M	Married	Low	2317.93
386.321667						
923	924	23	M	Single	High	760.13
380.065000						
924	925	54	F	Divorced	Medium	1930.41
386.082000						
989	990	37	F	Widowed	High	483.50
483.500000						

	NumTransactions	NumInteractions	MostCommonResolution	
LastLoginDate \				
13619	1	NaN	NaN	2023-06-
14318	1	1.0	Resolved	2023-02-
23613	1	NaN	NaN	2023-01-
26029	1	1.0	Unresolved	2023-05-
28122	5	1.0	Resolved	2023-07-
48714	1	NaN	NaN	2023-06-
49728	2	NaN	NaN	2023-06-
50904	1	1.0	Unresolved	2023-07-
56422	6	2.0	Resolved	2023-04-
61325	3	1.0	Resolved	2023-03-
62917	2	1.0	Unresolved	2023-08-
82609	6	NaN	NaN	2023-10-
92325	2	1.0	Unresolved	2023-03-
92420	5	NaN	NaN	2023-01-
98907	1	2.0	Resolved	2023-06-

	LoginFrequency	ServiceUsage	ChurnStatus
136	38	Online Banking	1
143	39	Online Banking	1
236	16	Online Banking	1

260	49	Online Banking	1
281	3	Online Banking	1
487	33	Mobile App	1
497	5	Website	1
509	31	Mobile App	1
564	30	Website	1
613	24	Mobile App	1
629	44	Website	1
826	38	Online Banking	1
923	41	Website	1
924	38	Mobile App	1
989	49	Mobile App	1

□ AvgSpent vs. ChurnStatus

The distribution of AvgSpent reveals a **slightly right-skewed curve** with a central cluster between \$150–\$300. When split by ChurnStatus, both churned and retained customers show similar spread, but a **slightly higher proportion of churners exists in the upper-middle spend range**.

□ Insight:

- High spend does **not guarantee loyalty**. A subset of churners are high spenders, potentially indicating dissatisfaction or lack of personalization despite high engagement.
- This segment should be considered for **targeted retention strategies**, such as loyalty rewards, concierge support, or exclusive financial products.